

pblas_wrapper

NAME

pblas_wrapper - a frovedis module provides user-friendly interfaces for commonly used pblas routines in scientific applications like machine learning algorithms.

SYNOPSIS

```
#include <frovedis/matrix/pblas_wrapper.hpp>
```

WRAPPER FUNCTIONS

```
void swap (const sliced_blockcyclic_vector<T>& v1,
           const sliced_blockcyclic_vector<T>& v2)
void copy (const sliced_blockcyclic_vector<T>& v1,
           const sliced_blockcyclic_vector<T>& v2)
void scal (const sliced_blockcyclic_vector<T>& v,
           T al)
void axpy (const sliced_blockcyclic_vector<T>& v1,
           const sliced_blockcyclic_vector<T>& v2,
           T al = 1.0)
T dot (const sliced_blockcyclic_vector<T>& v1,
       const sliced_blockcyclic_vector<T>& v2)
T nrm2 (const sliced_blockcyclic_vector<T>& v)
void gemv (const sliced_blockcyclic_matrix<T>& m,
           const sliced_blockcyclic_vector<T>& v1,
           const sliced_blockcyclic_vector<T>& v2,
           char trans = 'N',
           T al = 1.0,
           T be = 0.0)
void ger (const sliced_blockcyclic_vector<T>& v1,
          const sliced_blockcyclic_vector<T>& v2,
          const sliced_blockcyclic_matrix<T>& m,
          T al = 1.0)
void gemm (const sliced_blockcyclic_matrix<T>& m1,
           const sliced_blockcyclic_matrix<T>& m2,
           const sliced_blockcyclic_matrix<T>& m3,
           char trans_m1 = 'N',
           char trans_m2 = 'N',
           T al = 1.0,
           T be = 0.0)
```

```

void geadd (const sliced_blockcyclic_matrix<T>& m1,
           const sliced_blockcyclic_matrix<T>& m2,
           char trans = 'N',
           T al = 1.0,
           T be = 1.0)

```

OVERLOADED OPERATORS

```

blockcyclic_matrix<T>
operator* (const sliced_blockcyclic_matrix<T>& m1,
          const sliced_blockcyclic_matrix<T>& m2)
blockcyclic_matrix<T>
operator+ (const sliced_blockcyclic_matrix<T>& m1,
          const sliced_blockcyclic_matrix<T>& m2)
blockcyclic_matrix<T>
operator~ (const sliced_blockcyclic_matrix<T>& m1)

```

DESCRIPTION

PBLAS is a high-performance external library written in Fortran language. It provides rich set of functionalities on vectors and matrices. The computation loads of these functionalities are parallelized over the available processes in a system and the user interfaces of this library is very detailed and complex in nature. It requires a strong understanding on each of the input parameters, along with some distribution concepts.

Frovedis provides a wrapper module for some commonly used PBLAS subroutines in scientific applications like machine learning algorithms. These wrapper interfaces are very simple and user needs not to consider all the detailed distribution parameters. Only specifying the target vectors or matrices with some other parameters (depending upon need) are fine. At the same time, all the use cases of a PBLAS routine can also be performed using Frovedis PBLAS wrapper of that routine.

These wrapper routines are global functions in nature. Thus they can be called easily from within the “frovedis” namespace. As a distributed input vector, they accept “`blockcyclic_matrix<T>`” with single column or “`sliced_blockcyclic_vector<T>`”. And as a distributed input matrix, they accept “`blockcyclic_matrix<T>`” or “`sliced_blockcyclic_matrix<T>`”. “T” is a template type which can be either “float” or “double”. The individual detailed descriptions can be found in the subsequent sections. Please note that the term “inout”, used in the below section indicates a function argument as both “input” and “output”.

Detailed Description

swap (v1, v2)

Parameters

v1: A `blockcyclic_matrix<T>` with single column or a `sliced_blockcyclic_vector<T>` (inout)
v2: A `blockcyclic_matrix<T>` with single column or a `sliced_blockcyclic_vector<T>` (inout)

Purpose

It will swap the contents of v1 and v2, if they are semantically valid and are of same length.

Return Value

On success, it returns void. If any error occurs, it throws an exception.

copy (v1, v2)

Parameters

v1: A `blockcyclic_matrix<T>` with single column or a `sliced_blockcyclic_vector<T>` (input)

v2: A `blockcyclic_matrix<T>` with single column or a `sliced_blockcyclic_vector<T>` (output)

Purpose

It will copy the contents of *v1* in *v2* ($v2 = v1$), if they are semantically valid and are of same length.

Return Value

On success, it returns void. If any error occurs, it throws an exception.

scal (v, al)

Parameters

v: A `blockcyclic_matrix<T>` with single column or a `sliced_blockcyclic_vector<T>` (inout)

al: A “T” typed argument (float or double) to specify the value to which the input vector needs to be scaled. (input)

Purpose

It will scale the input vector with the provided “al” value, if it is semantically valid. On success, input vector “v” would be updated (in-place scaling).

Return Value

On success, it returns void. If any error occurs, it throws an exception.

axpy (v1, v2, al=1.0)

Parameters

v1: A `blockcyclic_matrix<T>` with single column or a `sliced_blockcyclic_vector<T>` (input)

v2: A `blockcyclic_matrix<T>` with single column or a `sliced_blockcyclic_vector<T>` (inout)

al: A “T” typed argument (float or double) to specify the value to which “v1” needs to be scaled (not in-place scaling) [Default: 1.0] (input/optional)

Purpose

It will solve the expression $v2 = al*v1 + v2$, if the input vectors are semantically valid and are of same length. On success, “v2” will be updated with desired result, but “v1” would remain unchanged.

Return Value

On success, it returns void. If any error occurs, it throws an exception.

dot (v1, v2)

Parameters

v1: A `blockcyclic_matrix<T>` with single column or a `sliced_blockcyclic_vector<T>` (input)

v2: A `blockcyclic_matrix<T>` with single column or a `sliced_blockcyclic_vector<T>` (input)

Purpose

It will perform dot product of the input vectors, if they are semantically valid and are of same length. Input vectors would not get modified during the operation.

Return Value

On success, it returns the dot product result of the type “float” or “double”. If any error occurs, it throws an exception.

nrm2 (v)

Parameters

v: A `blockcyclic_matrix<T>` with single column or a `sliced_blockcyclic_vector<T>` (input)

Purpose

It will calculate the norm of the input vector, if it is semantically valid. Input vector would not get modified during the operation.

Return Value

On success, it returns the norm value of the type “float” or “double”. If any error occurs, it throws an exception.

gemv (m, v1, v2, trans='N', al=1.0, be=0.0)

Parameters

m: A `blockcyclic_matrix<T>` or a `sliced_blockcyclic_matrix<T>` (input)

v1: A `blockcyclic_matrix<T>` with single column or a `sliced_blockcyclic_vector<T>` (input)

v2: A `blockcyclic_matrix<T>` with single column or a `sliced_blockcyclic_vector<T>` (inout)

trans: A character value can be either ‘N’ or ‘T’ [Default: ‘N’] (input/optional)

al: A “T” typed (float or double) scalar value [Default: 1.0] (input/optional)

be: A “T” typed (float or double) scalar value [Default: 0.0] (input/optional)

Purpose

The primary aim of this routine is to perform simple matrix-vector multiplication.

But it can also be used to perform any of the below operations:

- (1) $v2 = al * m * v1 + be * v2$
- (2) $v2 = al * \text{transpose}(m) * v1 + be * v2$

If *trans*='N', then expression (1) is solved. In that case, the size of “v1” must be at least the number of columns in “m” and the size of “v2” must be at least the number of rows in “m”.

If *trans*='T', then expression (2) is solved. In that case, the size of “v1” must be at least the number of rows in “m” and the size of “v2” must be at least the number of columns in “m”.

Since “v2” is used as input-output both, memory must be allocated for this vector before calling this routine, even if simple matrix-vector multiplication is required. Otherwise, this routine will throw an exception.

For simple matrix-vector multiplication, no need to specify values for the input parameters “trans”, “al” and “be” (leave them at their default values).

On success, “v2” will be overwritten with the desired output. But “m” and “v1” would remain unchanged.

Return Value

On success, it returns void. If any error occurs, it throws an exception.

ger (v1, v2, m, al=1.0)

Parameters

v1: A `blockcyclic_matrix<T>` with single column or a `sliced_blockcyclic_vector<T>` (input)

v2: A `blockcyclic_matrix<T>` with single column or a `sliced_blockcyclic_vector<T>` (input)

m: A `blockcyclic_matrix<T>` or a `sliced_blockcyclic_matrix<T>` (inout)

al: A “T” typed (float or double) scalar value [Default: 1.0] (input/optional)

Purpose

The primary aim of this routine is to perform simple vector-vector multiplication of the sizes “a” and “b” respectively to form an axb matrix. But it can also be used to perform the below operations:

```
m = al*v1*v2' + m
```

This operation can only be performed if the inputs are semantically valid and the size of “v1” is at least the number of rows in matrix “m” and the size of “v2” is at least the number of columns in matrix “m”.

Since “m” is used as input-output both, memory must be allocated for this matrix before calling this routine, even if simple vector-vector multiplication is required. Otherwise it will throw an exception.

For simple vector-vector multiplication, no need to specify the value for the input parameter “al” (leave it at its default value).

On success, “m” will be overwritten with the desired output. But “v1” and “v2” will remain unchanged.

Return Value

On success, it returns void. If any error occurs, it throws an exception.

gemm (m1, m2, m3, trans__m1='N', trans__m2='N', al=1.0, be=0.0)

Parameters

m1: A `blockcyclic_matrix<T>` or a `sliced_blockcyclic_matrix<T>` (input)
m2: A `blockcyclic_matrix<T>` or a `sliced_blockcyclic_matrix<T>` (input)
m3: A `blockcyclic_matrix<T>` or a `sliced_blockcyclic_matrix<T>` (inout)
trans__m1: A character value can be either 'N' or 'T' [Default: 'N'] (input/optional)
trans__m2: A character value can be either 'N' or 'T' [Default: 'N'] (input/optional)
al: A “T” typed (float or double) scalar value [Default: 1.0] (input/optional)
be: A “T” typed (float or double) scalar value [Default: 0.0] (input/optional)

Purpose

The primary aim of this routine is to perform simple matrix-matrix multiplication. But it can also be used to perform any of the below operations:

- (1) $m3 = al*m1*m2 + be*m3$
- (2) $m3 = al*transpose(m1)*m2 + be*m3$
- (3) $m3 = al*m1*transpose(m2) + be*m3$
- (4) $m3 = al*transpose(m1)*transpose(m2) + be*m3$

(1) will be performed, if both “trans__m1” and “trans__m2” are 'N'.

(2) will be performed, if trans__m1='T' and trans__m2 = 'N'.

(3) will be performed, if trans__m1='N' and trans__m2 = 'T'.

(4) will be performed, if both “trans__m1” and “trans__m2” are 'T'.

If we have four variables `nrowa`, `nrowb`, `ncola`, `ncolb` defined as follows:

```
if(trans__m1 == 'N') {
    nrowa = number of rows in m1
    ncola = number of columns in m1
}
else if(trans__m1 == 'T') {
    nrowa = number of columns in m1
    ncola = number of rows in m1
}
```

```

if(trans_m2 == 'N') {
    nrowb = number of rows in m2
    ncolb = number of columns in m2
}
else if(trans_m2 == 'T') {
    nrowb = number of columns in m2
    ncolb = number of rows in m2
}

```

Then this function can be executed successfully, if the below conditions are all true:

- (a) "ncola" is equal to "nrowb"
- (b) number of rows in "m3" is equal to or greater than "nrowa"
- (b) number of columns in "m3" is equal to or greater than "ncolb"

Since “m3” is used as input-output both, memory must be allocated for this matrix before calling this routine, even if simple matrix-matrix multiplication is required. Otherwise it will throw an exception.

For simple matrix-matrix multiplication, no need to specify the value for the input parameters “trans_m1”, “trans_m2”, “al”, “be” (leave them at their default values).

On success, “m3” will be overwritten with the desired output. But “m1” and “m2” will remain unchanged.

Return Value

On success, it returns void. If any error occurs, it throws an exception.

geadd (m1, m2, trans='N', al=1.0, be=1.0)

Parameters

m1: A `blockcyclic_matrix<T>` or a `sliced_blockcyclic_matrix<T>` (input)
m2: A `blockcyclic_matrix<T>` or a `sliced_blockcyclic_matrix<T>` (inout)
trans: A character value can be either 'N' or 'T' [Default: 'N'] (input/optional)
al: A “T” typed (float or double) scalar value [Default: 1.0] (input/optional)
be: A “T” typed (float or double) scalar value [Default: 1.0] (input/optional)

Purpose

The primary aim of this routine is to perform simple matrix-matrix addition. But it can also be used to perform any of the below operations:

- (1) $m2 = al*m1 + be*m2$
- (2) $m2 = al*transpose(m1) + be*m2$

If trans='N', then expression (1) is solved. In that case, the number of rows and the number of columns in “m1” should be equal to the number of rows and the number of columns in “m2” respectively.

If trans='T', then expression (2) is solved. In that case, the number of columns and the number of rows in “m1” should be equal to the number of rows and the number of columns in “m2” respectively.

If it is needed to scale the input matrices before the addition, corresponding “al” and “be” values can be provided. But for simple matrix-matrix addition, no need to specify values for the input parameters “trans”, “al” and “be” (leave them at their default values).

On success, “m2” will be overwritten with the desired output. But “m1” would remain unchanged.

Return Value

On success, it returns void. If any error occurs, it throws an exception.

operator* (m1, m2)

Parameters

m1: A `blockcyclic_matrix<T>` or a `sliced_blockcyclic_matrix<T>` (input)

m2: A `blockcyclic_matrix<T>` or a `sliced_blockcyclic_matrix<T>` (input)

Purpose

This operator operates on two input matrices and returns the resultant matrix after successful multiplication. Both the input matrices remain unchanged.

Return Value

On success, it returns the resultant matrix of the type `blockcyclic_matrix<T>`. If any error occurs, it throws an exception.

operator+ (m1, m2)

Parameters

m1: A `blockcyclic_matrix<T>` or a `sliced_blockcyclic_matrix<T>` (input)

m2: A `blockcyclic_matrix<T>` or a `sliced_blockcyclic_matrix<T>` (input)

Purpose

This operator operates on two input matrices and returns the resultant matrix after successful addition. Both the input matrices remain unchanged.

Return Value

On success, it returns the resultant matrix of the type `blockcyclic_matrix<T>`. If any error occurs, it throws an exception.

operator~ (m1)

Parameters

m1: A `blockcyclic_matrix<T>` or a `sliced_blockcyclic_matrix<T>` (input)

Purpose

This operator operates on single input matrix and returns its transposed matrix. E.g., if “m” is a matrix of type `blockcyclic_matrix<T>`, then “~m” will return transposed of matrix “m” of the type `blockcyclic_matrix<T>`.

Return Value On success, it returns the resultant matrix of the type `blockcyclic_matrix<T>`. If any error occurs, it throws an exception.

SEE ALSO

`sliced_blockcyclic_matrix_local`, `sliced_blockcyclic_vector_local`, `blas_wrapper`