

Ridge Regression

NAME

Ridge Regression - A regression algorithm supported by Frowedis to predict the continuous output with L2 regularization.

SYNOPSIS

```
#include <frowedis/ml/glm/ridge_regression_with_sgd.hpp>

linear_regression_model<T>
ridge_regression_with_sgd::train (crs_matrix<T>& data,
    dvector<T>& label,
    size_t numIteration = 1000,
    T alpha = 0.01,
    T miniBatchFraction = 1.0,
    T regParam = 0.01,
    bool isIntercept = false,
    T convergenceTol = 0.001,
    MatType mType = HYBRID)

linear_regression_model<T>
ridge_regression_with_sgd::train (crs_matrix<T>& data,
    dvector<T>& label,
    linear_regression_model<T>& initModel,
    size_t numIteration = 1000,
    T alpha = 0.01,
    T miniBatchFraction = 1.0,
    T regParam = 0.01,
    bool isIntercept = false,
    T convergenceTol = 0.001,
    MatType mType = HYBRID)

#include <frowedis/ml/glm/ridge_regression_with_lbfgs.hpp>

linear_regression_model<T>
ridge_regression_with_lbfgs::train (crs_matrix<T>& data,
    dvector<T>& label,
    size_t numIteration = 1000,
    T alpha = 0.01,
    size_t hist_size = 10,
    T regParam = 0.01,
    bool isIntercept = false,
    T convergenceTol = 0.001,
    MatType mType = HYBRID)
```

```

linear_regression_model<T>
ridge_regression_with_lbfgs::train (crs_matrix<T>& data,
    dvector<T>& label,
    linear_regression_model<T>& initModel,
    size_t numIteration = 1000,
    T alpha = 0.01,
    size_t hist_size = 10,
    T regParam = 0.01,
    bool isIntercept = false,
    T convergenceTol = 0.001,
    MatType mType = HYBRID)

```

DESCRIPTION

Linear least squares is the most common formulation for regression problems. It is a linear method with the loss function given by the **squared loss**:

$$L(\mathbf{w}; \mathbf{x}, y) := 1/2(\mathbf{w}^T \mathbf{x} - y)^2$$

Where the vectors \mathbf{x} are the training data examples and y are their corresponding labels which we want to predict. \mathbf{w} is the linear model (also known as weight) which uses a single weighted sum of features to make a prediction. The method is called linear since it can be expressed as a function of $\mathbf{w}^T \mathbf{x}$ and y . Ridge regression uses L2 regularization to address the overfit problem.

The gradient of the squared loss is: $(\mathbf{w}^T \mathbf{x} - y) \cdot \mathbf{x}$

The gradient of the regularizer is: \mathbf{w}

Frovedis provides implementation of ridge regression with two different optimizers: (1) stochastic gradient descent with minibatch and (2) LBFGS optimizer.

The simplest method to solve optimization problems of the form **min** $f(\mathbf{w})$ is gradient descent. Such first-order optimization methods well-suited for large-scale and distributed computation. Whereas, L-BFGS is an optimization algorithm in the family of quasi-Newton methods to solve the optimization problems of the similar form.

Like the original BFGS, L-BFGS (Limited Memory BFGS) uses an estimation to the inverse Hessian matrix to steer its search through feature space, but where BFGS stores a dense $n \times n$ approximation to the inverse Hessian (n being the number of features in the problem), L-BFGS stores only a few vectors that represent the approximation implicitly. L-BFGS often achieves rapider convergence compared with other first-order optimization.

Detailed Description

ridge_regression_with_sgd::train()

Parameters

data: A `crs_matrix<T>` containing the sparse feature matrix

label: A `dvector<T>` containing the output labels

numIteration: A `size_t` parameter containing the maximum number of iteration count (Default: 1000)

alpha: A parameter of `T` type containing the learning rate (Default: 0.01)

minibatchFraction: A parameter of `T` type containing the minibatch fraction (Default: 1.0)

regParam: A parameter of `T` type containing the regularization parameter (also known as λ) (Default: 0.01)

isIntercept: A boolean parameter to specify whether to include intercept term (bias term) or not (Default: false)
convergenceTol: A parameter of T type containing the threshold value to determine the convergence (Default: 0.001)
mType: `frovedis::MatType` parameter specifying the matrix type to be used for internal calculation (Default: HYBRID for SX architecture, CRS for other architectures)

Purpose

It trains a linear regression model with stochastic gradient descent with minibatch optimizer and with L2 regularization. It starts with an initial guess of zeros for the model vector and keeps updating the model to minimize the cost function until convergence is achieved or maximum iteration count is reached. After the training, it returns the trained output model.

Return Value

After the successful training, it returns a trained model of the type `linear_regression_model<T>`.

`ridge_regression_with_sgd::train()`

Parameters

data: A `crs_matrix<T>` containing the sparse feature matrix
label: A `dvector<T>` containing the output labels
initModel: A `linear_regression_model<T>` containing the user provided initial model values
numIteration: A `size_t` parameter containing the maximum number of iteration count (Default: 1000)
alpha: A parameter of T type containing the learning rate (Default: 0.01)
minibatchFraction: A parameter of T type containing the minibatch fraction (Default: 1.0)
regParam: A parameter of T type containing the regularization parameter (also known as lambda) (Default: 0.01)
isIntercept: A boolean parameter to specify whether to include intercept term (bias term) or not (Default: false)
convergenceTol: A parameter of T type containing the threshold value to determine the convergence (Default: 0.001)
mType: `frovedis::MatType` parameter specifying the matrix type to be used for internal calculation (Default: HYBRID for SX architecture, CRS for other architectures)

Purpose

It trains a linear regression model with stochastic gradient descent with minibatch optimizer and with L2 regularization. Instead of an initial guess of zeors, it starts with user provided initial model values and keeps updating the model to minimize the cost function until convergence is achieved or maximum iteration count is reached. After the training, it returns the trained output model.

Return Value

After the successful training, it returns a trained model of the type `linear_regression_model<T>`.

`ridge_regression_with_lbfgs::train()`

Parameters

data: A `crs_matrix<T>` containing the sparse feature matrix
label: A `dvector<T>` containing the output labels
numIteration: A `size_t` parameter containing the maximum number of iteration count (Default: 1000)
alpha: A parameter of T type containing the learning rate (Default: 0.01)
hist_size: A parameter of `size_t` type containing the number of gradient history to be stored (Default: 10)
regParam: A parameter of T type containing the regularization parameter (also known as lambda) (Default: 0.01)
isIntercept: A boolean parameter to specify whether to include intercept term (bias term) or not (Default: false)

convergenceTol: A parameter of T type containing the threshold value to determine the convergence (Default: 0.001)

mType: `frovedis::MatType` parameter specifying the matrix type to be used for internal calculation (Default: HYBRID for SX architecture, CRS for other architectures)

Purpose

It trains a linear regression model with LBFGS optimizer and with L2 regularization. It starts with an initial guess of zeros for the model vector and keeps updating the model to minimize the cost function until convergence is achieved or maximum iteration count is reached. After the training, it returns the trained output model.

Return Value

After the successful training, it returns a trained model of the type `linear_regression_model<T>`.

`ridge_regression_with_lbfgs::train()`

Parameters

data: A `crs_matrix<T>` containing the sparse feature matrix

label: A `dvector<T>` containing the output labels

initModel: A `linear_regression_model<T>` containing the user provided initial model values

numIteration: A `size_t` parameter containing the maximum number of iteration count (Default: 1000)

alpha: A parameter of T type containing the learning rate (Default: 0.01)

hist_size: A parameter of `size_t` type containing the number of gradient history to be stored (Default: 10)

regParam: A parameter of T type containing the regularization parameter (also known as lambda) (Default: 0.01)

isIntercept: A boolean parameter to specify whether to include intercept term (bias term) or not (Default: false)

convergenceTol: A parameter of T type containing the threshold value to determine the convergence (Default: 0.001)

mType: `frovedis::MatType` parameter specifying the matrix type to be used for internal calculation (Default: HYBRID for SX architecture, CRS for other architectures)

Purpose

It trains a linear regression model with LBFGS optimizer and with L2 regularizer. Instead of an initial guess of zeors, it starts with user provided initial model values and keeps updating the model to minimize the cost function until convergence is achieved or maximum iteration count is reached. After the training, it returns the trained output model.

Return Value

After the successful training, it returns a trained model of the type `linear_regression_model<T>`.

SEE ALSO

`linear_regression_model`, `linear_regression`, `lasso_regression`